

# OpenMP mini-app practical

Vasileios Karakasis

Swiss National Supercomputing Center (CSCS)

*Slides prepared by Ben Cumming*



Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

CSCS

Swiss National Supercomputing Centre



# Adding OpenMP to the mini-app

- The aim of this practical is to take the serial mini-app implementation, and make it faster with OpenMP – while still getting the correct answer!

# Hints

- Before starting find two sets of input parameters that converge for the serial version
  - note the time to solution
    - you want this to get faster as you add OpenMP
    - **NOTE:** but you might have to add quite a few directives before things actually get faster
  - note the number of conjugate gradient iterations
    - use this to check after you add each directive that you are still getting the right answer
    - **NOTE:** remember that there will be some small variations because floating point operations are not commutative



# First test

- Get the code, by checking it out from github

```
> git pull
> cd <SummerSchool2017path>/miniapp/openmp
> ls
cxx fortran
> cd cxx
> make
> srun ./main 128 128 100 0.01
<note time to solution and conjugate gradient iterations>
> srun ./main 256 256 200 0.01
<note time to solution and conjugate gradient iterations>
```

I choose the C++ version here

# Step 1

- replace the welcome message in main.cpp/main.f90 with a message that tells the user
  - that this is the openmp version
  - how many threads it is using

```
> make
> OMP_NUM_THREADS=12 srun -n1 -c12 -hint=nomulithread ./main 256 256 200 0.01
...
=====
                welcome to mini-stencil!  version ::
Fortran90 openMP with 12 threads
```

# Step 2: Linear Algebra

- Open [linalg.cpp/f90](#) and add directives to the functions subroutines `ss_XXXX`
  - do one or two at a time
    - recompile frequently and run with 12 threads to check that you are still getting the right answer
- Once finished with that file, did your changes make any improvement?
  - compare the 128x128 and 256x256 results



# Step 3: the diffusion stencil

- The final step is to parallelize the stencil operator in `operators.cpp/f90`.
- The nested for/do loop is an obvious target
  - it covers  $n_x \times n_y$  grid points
- How about the boundary loops?



# Step 4: testing

- how does it scale at different resolutions?
  - 32x32
  - 64x64
  - 128x128
  - 256x256
  - 512x512
  - 1024x1024
- Advanced C++:
  - can you implement first touch memory allocation in the C++ version?
  - requires adding just one OpenMP directive





# Extras

- can you implement first touch memory allocation
  - requires adding just one OpenMP directive in C++
- does the stencil kernel vectorize?
  - look at Cray and Intel compiler vectorization reports
  - can you make it vectorize?



# Thanks for your attention



Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

CSCS

Swiss National Supercomputing Centre

